PyTupli: Enabling Collaboration in Offline Reinforcement Learning

10.12.2025

Summary

Offline reinforcement learning (RL) offers a powerful way to derive effective decision-making policies for control problems using pre-collected data. However, managing and sharing such datasets in collaborative research settings can be challenging, as proper versioning, filtering, and access control are required. PyTupli is a free, Python-based toolkit designed to support this workflow by providing an easy-to-use yet specialized framework that remains independent of external service providers. Unlike centralized platforms, PyTupli is designed for self-hosted deployment, giving research teams full control over their data infrastructure and access policies. Its client library allows users to serialize and store control problems, upload new data, and retrieve precisely the subsets they need through flexible and expressive filters. Built-in metrics help evaluate dataset coverage and utility, informing both dataset selection and algorithm design for offline RL. To ensure secure deployment, a container-based server component offers authentication, role-based access control, and automated certificate provisioning. Together, these capabilities enable researchers to create, manage, exchange, and analyze datasets for offline RL in a robust and accessible manner.

Statement of need

Reinforcement learning (RL) provides powerful methods for decision-making under uncertainty, but training RL agents typically requires extensive interaction with the underlying system or a computationally expensive simulator. Offline RL has emerged as a paradigm that alleviates this requirement by training agents solely on previously collected data (Lange, Gabel, and Riedmiller 2012). Such datasets contain tuples consisting of *state, action, next state,* and *reward*, obtained from recordings of real systems or generated through simulators.

Effectively managing, sharing, and curating these datasets is essential for collaborative offline RL research, yet existing tools provide only partial support. Platforms like Zenodo or the free version of HuggingFace allow users to share

finalized datasets with a general public but are not suitable for ongoing or private collaborations. Furthermore, they lack mechanisms for tracking internal dataset structure or performing efficient, fine-grained queries. The same applies to version control systems such as GitHub. Traditional databases are better suited for this purpose but require substantial expertise to design and maintain robust workflows.

PyTupli addresses this gap by providing a dedicated Python toolkit for creating, storing, and sharing tuple datasets for custom environments. We provide a Docker container that each research group or collaboration can deploy independently to maintain complete control over their data. Each dataset is associated to a benchmark. Benchmarks are stored as JSON-serialized objects and can be linked to related artifacts, including time-series data, algorithm hyperparameters, or trained policies, allowing multiple benchmarks to reference shared resources. Because it is built for scalable collaboration, PyTupli includes integrated user and access management features.

Since the performance of offline RL algorithms often depends critically on the quality of the underlying dataset (Schweighofer et al. 2022; Suttle, Suresh, and Nieto-Granda 2025; Asadulaev, Karray, and Takac 2025), PyTupli offers extensive filtering capabilities. Whereas established offline RL datasets primarily support filtering by entire episodes (Younis et al. 2024; Liu et al. 2023), PyTupli enables tuple-level filtering as well. This can be used, for example, to rebalance datasets with sparse rewards or selectively include transitions from specific regions of the state space. In addition, PyTupli provides a suite of metrics for assessing dataset coverage and reward characteristics, which can serve as predictors of offline RL performance (Schweighofer et al. 2022; Asadulaev, Karray, and Takac 2025) and help guide algorithm selection.

State of the Field

Publicly available tuple datasets have been essential for advancing offline RL algorithms (Kumar et al. 2020; Kostrikov, Nair, and Levine 2021). These curated collections span various domains, such as robotics and games (Fu et al. 2020; Younis et al. 2024; Formanek et al. 2023; Gulcehre et al. 2020), power system control (Qin et al. 2022), and autonomous driving (Liu et al. 2023; Lee, Eom, and Kwon 2024), and are typically designed to support the development of improved offline RL methods. As offline RL techniques mature, they are being applied to increasingly diverse and task-specific control problems. Yet, to the best of our knowledge, no existing toolbox supports the collaborative creation, management, and sharing of datasets for custom environments. Minari (Younis et al. 2024) is the closest related tool, offering a repository of standardized datasets along with functionality for filtering, environment reconstruction, and recording new interactions. Its focus, however, remains on distributing datasets for established benchmarks.

PyTupli instead targets collaborative workflows for custom control tasks. It enables researchers to share datasets and benchmarks directly within project teams without dependence on a central public server. Although Minari permits users to request publication on its official platform, this approach is often unsuitable for ongoing or proprietary work or for datasets that evolve over time. In addition, Minari does not provide tools for assessing dataset quality or coverage, which PyTupli includes to support informed dataset curation and algorithm selection in offline RL research.

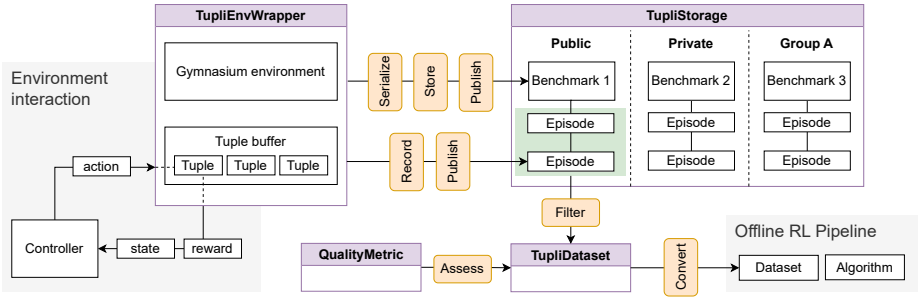

Figure 1: Overview of the core functionalities of PyTupli.

Core Functionalities

We briefly describe the core functionalities of PyTupli which are illustrated in Figure 1.

Benchmark and Artifact Management: PyTupli enables users to store any control task defined as a gymnasium environment. Environments may include parameterizable configurations that produce variations in task dynamics. A fully specified environment is stored as a benchmark, providing a unique reference for reproducible evaluation of controllers. Benchmarks can reference additional data, such as exogenous inputs, time-series data, or pre-trained models. These external units, referred to as artifacts, are stored independently and linked to multiple benchmarks to avoid duplication.

Data Management: PyTupli supports ingesting, storing, and querying structured datasets (RL tuples), including their relation to existing benchmark problems and any relevant metadata. MongoDB serves as the backend, providing scalable storage and efficient retrieval for large datasets. Table 1 shows how ingestion and retrieval times scale with dataset size.

Multi-User Collaboration and Access Control: PyTupli facilitates collaborative workflows through private, group, and public scopes. Based on their assigned role, users can store, retrieve, delete, and publish objects. A server-side backend with FastAPI provides a REST interface for secure, programmatic

access, while token-based authentication ensures secure sharing across teams or organizations.

Integration with Existing Offline RL Infrastructure: An interface to the gymnasium framework enables users to record interactions with gymnasium environments as RL tuples. Furthermore, retrieved tuple datasets are made available in a form that can easily be converted into the dataset formats used by existing offline RL libraries such as d3rlpy (Seno and Imai 2022).

Assessment of Dataset Quality: PyTupli implements metrics to evaluate dataset quality in terms of coverage and expected returns. These metrics inform dataset selection and can provide guidance for algorithm choice. Detailed formulations are provided in the following section.

Table 1: Upload and download times for established datasets averaged over 10 runs. We chose two examples with low, medium, and high dataset size from the Minari collection. However, not only the size, but also the nature of observations has a strong influence on processing times.

Dataset	Size	M	N	Upload (s)	Download (s)
D4RL					
door/human-v2	3.5MB	25	7K	0.83	0.24
hammer/human-v2	6.2MB	25	11K	1.11	0.40
antmaze/medium-play-v1	605.2MB	1K	1M	173.26	126.62
Atari					
pitfall/expert-v0	351.7MB	10	65K	18.56	16.51
Mujoco					
ant/expert-v0	1.92GB	2K	2M	64.17	29.65
humanoid/expert-v0	2.95GB	1K	999K	96.61	55.32

Quality Metrics

Return-Based Metrics

Return-based metrics, such as trajectory quality (TQ) (Schweighofer et al. 2022) or the average Q-value (Asadulaev, Karray, and Takac 2025) can inform algorithm decision. For example, Schweighofer et al. (2022) show that behavioral cloning performs well despite its simplicity if the dataset has a high TQ. For datasets with low TQ, algorithms from the deep Q-network family perform well in their experiments as they do not constrain the learned policy towards the distribution of the behavioral policy. TQ normalizes the average return of a dataset with respect to the returns obtained by a minimal performant and an expert policy. To provide similar insights without relying on such additional information, estimated relative return improvement (Swazinna, Udluft, and Runkler 2021) relates the

maximum trajectory return in the dataset to its average return. While estimated return improvement and TQ operate on a trajectory level, average Q-value estimation offers insights on the tuple level, making it a better predictor of offline RL performance (Asadulaev, Karray, and Takac 2025). It requires fitting a Q-function using Bellman updates, which is closely related to the objectives used in offline RL training. However, for continuous action spaces, the user needs to provide an evaluation policy to predict the next actions in the Bellman target. Such a policy can, for example, be obtained using behavioral cloning on the dataset.

Coverage-Based Metrics

An important question when assessing a dataset is whether the behavioral policy (or policies) used to generate it did explore the state and action space well enough to learn a meaningful target policy from the data. A common approach for quantifying explorativeness is to approximate the entropy of transition probabilities for the behavior policy. For discrete state and action spaces, Schweighofer et al. (2022) suggest to approximate this by counting unique state-action pairs. Optionally, this value can be normalized using a reference dataset $\mathcal{D}_{\text{ref}}$ of same size, for example, the replay buffer collected during online training. Schweighofer et al. (2022) show that low state-action coverage values hinder performance of a large variety of algorithms. While counting unique state-action pairs aims at estimating the Shannon entropy of the transition probabilities, Suttle, Suresh, and Nieto-Granda (2025) suggest that datasets that maximize their proposed behavioral entropy metric support better offline RL performance. They suggest a k-nearest-neighbor estimator of the true behavioral entropy that relies on density-based weighting of different regions in the state-action space.

Software Design

PyTupli is designed around a clear separation between a lightweight client library and a centralized server backend. On the client side, PyTupli integrates via a Gymnasium environment wrapper, which is a well-established abstraction in the RL ecosystem. The server component exposes a REST API that implements functionality specific to offline RL datasets, such as structured storage of benchmarks, episodes, and artifacts, as well as flexible filtering and access control. A centralized service was favored over object storage or git-based workflows because offline RL datasets require domain-specific querying and metadata handling that these approaches do not natively support. Deployment is simplified via a Docker Compose setup, providing a production-ready stack that can be launched with minimal configuration. Here, the API server is hidden behind an Nginx web server acting as a reverse proxy for increased scalability. We chose MongoDB as the backend due to its ability to store heterogeneous, evolving data without rigid schemas while supporting efficient indexing for nested fields. GridFS enables integrated storage of large artifacts, avoiding external dependencies.

Existing tools such as Minari focus on distributing pre-existing datasets for offline RL. PyTupli instead targets research groups that need to create, host, and curate custom benchmarks collaboratively. This fundamental difference in scope and architecture made contributing to existing projects impractical, motivating the development of new software tailored to this use case.

Research Impact Statement

Since its initial release on PyPI in May 2025, PyTupli has been downloaded approximately 10,000 times, suggesting early uptake by practitioners and researchers. The project is actively maintained by a small but committed team of two contributors, with ongoing development supported by a full automated test suite and a CI/CD pipeline that ensures reliability and code quality.

Although PyTupli has not yet accumulated extensive academic citations, it demonstrates strong indicators of potential research impact through its technical maturity and readiness for community adoption. Comprehensive documentation is publicly available via ReadTheDocs (https://pytupli.readthedocs.io/en/latest/), lowering the barrier to entry for new users. In addition, a set of tutorials (Jupyter notebooks) provides practical, hands-on guidance, enabling users to quickly integrate PyTupli into their own workflows.

Importantly, PyTupli has already been integrated into external research software, most notably the CommonPower framework for reinforcement learning–based power system control. This integration demonstrates PyTupli's applicability in real research settings and its potential to support future work in applied offline reinforcement learning. Together, these elements indicate that PyTupli is a robust, reusable, and community-ready research tool with growing impact potential.

AI Usage Disclosure

Visual Studio Code with Claude Sonnet 4.5 was used for the scaffolding of the test in PyTupli. GPT 5.2 was used for refining the present manuscript.

Acknowledgements

This work was partially supported by the German Research Foundation (AL 1185/9-1).

References

Asadulaev, Arip, Fakhri Karray, and Martin Takac. 2025. "Expert or Not? Assessing Data Quality in Offline Reinforcement Learning." *arXiv Preprint arXiv:2510.12638*. https://doi.org/10.48550/arXiv.2510.12638.

Formanek, Claude, Asad Jeewa, Jonathan Shock, and Arnu Pretorius. 2023. "Off-the-Grid MARL: Datasets and Baselines for Offline Multi-Agent Reinforcement Learning," 2442–44. https://doi.org/10.5555/3545946.3598961.

Fu, Justin, Aviral Kumar, Ofir Nachum, George Tucker, and Sergey Levine. 2020. "D4rl: Datasets for Deep Data-Driven Reinforcement Learning." *arXiv Preprint arXiv:2004.07219*. https://doi.org/https://doi.org/10.48550/arXiv.2004.07219.

Gulcehre, Caglar, Ziyu Wang, Alexander Novikov, Thomas Paine, Sergio Gómez, Konrad Zolna, Rishabh Agarwal, et al. 2020. "Rl Unplugged: A Suite of Benchmarks for Offline Reinforcement Learning." *Advances in Neural Information Processing Systems* 33: 7248–59. https://doi.org/10.5555/3495724.3496332.

Kostrikov, Ilya, Ashvin Nair, and Sergey Levine. 2021. "Offline Reinforcement Learning with Implicit q-Learning." *arXiv Preprint arXiv:2110.06169*. https://doi.org/10.48550/arXiv.2110.06169.

Kumar, Aviral, Aurick Zhou, George Tucker, and Sergey Levine. 2020. "Conservative q-Learning for Offline Reinforcement Learning." *Advances in Neural Information Processing Systems* 33: 1179–91. https://doi.org/10.5555/3495724.3495824.

Lange, Sascha, Thomas Gabel, and Martin Riedmiller. 2012. "Batch Reinforcement Learning." In *Reinforcement Learning: State-of-the-Art*, 45–73. Springer. https://doi.org/10.1007/978-3-642-27645-3_2.

Lee, Dongsu, Chanin Eom, and Minhae Kwon. 2024. "AD4RL: Autonomous Driving Benchmarks for Offline Reinforcement Learning with Value-Based Dataset." In *2024 IEEE International Conference on Robotics and Automation (ICRA)*, 8239–45. IEEE. https://doi.org/10.1109/ICRA57147.2024.10610308.

Liu, Zuxin, Zijian Guo, Haohong Lin, Yihang Yao, Jiacheng Zhu, Zhepeng Cen, Hanjiang Hu, et al. 2023. "Datasets and Benchmarks for Offline Safe Reinforcement Learning." *arXiv Preprint arXiv:2306.09303*. https://doi.org/10.48550/arXiv.2306.09303.

Qin, Rong-Jun, Xingyuan Zhang, Songyi Gao, Xiong-Hui Chen, Zewen Li, Weinan Zhang, and Yang Yu. 2022. "NeoRL: A Near Real-World Benchmark for Offline Reinforcement Learning." *Advances in Neural Information Processing Systems* 35: 24753–65. https://doi.org/10.5555/3600270.3602065.

Schweighofer, Kajetan, Marius-constantin Dinu, Andreas Radler, Markus Hofmarcher, Vihang Prakash Patil, Angela Bitto-Nemling, Hamid Eghbal-Zadeh, and Sepp Hochreiter. 2022. "A Dataset Perspective on Offline Reinforcement Learning." In *Conference on Lifelong Learning Agents*, 470–517. PMLR. https://doi.org/10.48550/arXiv.2111.04714.

Seno, Takuma, and Michita Imai. 2022. "D3rlpy: An Offline Deep Reinforcement Learning Library." *Journal of Machine Learning Research* 23 (315): 1–20. https://doi.org/10.5555/3586589.3586904.

Suttle, Wesley A, Aamodh Suresh, and Carlos Nieto-Granda. 2025. "Behavioral Entropy-Guided Dataset Generation for Offline Reinforcement Learning." *arXiv Preprint arXiv:2502.04141*. https://doi.org/10.48550/arXiv.2502.04141.

Swazinna, Phillip, Steffen Udluft, and Thomas Runkler. 2021. "Measuring Data Quality for Dataset Selection in Offline Reinforcement Learning." In *2021 IEEE Symposium Series on Computational Intelligence (SSCI)*, 1–8. IEEE. https://doi.org/10.1109/SSCI50451.2021.9660006.

Younis, Omar G., Rodrigo Perez-Vicente, John U. Balis, Will Dudley, Alex Davey, and Jordan K Terry. 2024. *Minari* (version 0.5.0). Zenodo. https://doi.org/10.5281/zenodo.13767625.